

### Problem statement 1:

Autonomous vehicles (AV) and intelligent transport systems (ITS) are the future of road transport. Automatic detection of vehicles on the road in real-time helps AV technology and makes ITS more intelligent in terms of vehicle tracking, vehicle counting, and road incident response.

### Objective part 1:

As the first part of this project, you need to develop an AI model using a deep learning framework that predicts the type of vehicle present in an image as well as localizes the vehicle by rectangular bounding box.

```
In [ ]: #Import the packages
from ultralytics import YOLO
import pandas as pd
import os
import shutil
from PIL import Image
```

### Object detection

1. Create a parent folder for custom model training and child folders to store data
2. Prepare the dataset for model training and keep the following points in mind while preparing it • This dataset contains many images, and depending on the compute power of the VM, it might take a very long time to unzip this huge amount of data

```
In [ ]: # Define paths
from ultralytics import YOLO
import pandas as pd
import os
import shutil
from PIL import Image

csv_path = "./sample_data/Part1/labels.csv"
img_folder = "./sample_data/Part1/images/"
yolo_dataset_path = "./sample_data/yolo_dataset/"
yolo_images_dir = os.path.join(yolo_dataset_path, "images")
yolo_labels_dir = os.path.join(yolo_dataset_path, "labels")

# Create necessary folders
for split in ['train', 'val']:
    os.makedirs(os.path.join(yolo_images_dir, split), exist_ok=True)
    os.makedirs(os.path.join(yolo_labels_dir, split), exist_ok=True)

# Load CSV
df = pd.read_csv(csv_path, header=None)
df.columns = ['image_id', 'class_name', 'x_min', 'y_min', 'x_max', 'y_max']

# Class name to ID mapping
class_map = {name: idx for idx, name in enumerate(df['class_name'].unique())}

# Group by image
grouped = df.groupby('image_id')
print(grouped.head())
# Simple train/val split
image_ids = df['image_id'].unique()
threshold = 5657
filtered_img = [s for s in image_ids if int(s) <= threshold]
split_idx = int(len(filtered_img) * 0.8)
train_ids = set(filtered_img[:split_idx])

for image_id, group in grouped:
    file_name = f"{str(image_id).zfill(8)}.jpg"
    orig_path = os.path.join(img_folder, file_name)

    if not os.path.exists(orig_path):
        continue # skip if image is missing

    # Determine split
    split = 'train' if image_id in train_ids else 'val'

    # Copy image
    dst_img_path = os.path.join(yolo_images_dir, split, file_name)
    shutil.copy(orig_path, dst_img_path)

    # Get image size
    with Image.open(orig_path) as img:
        width, height = img.size
```

```
# Convert annotations
yolo_lines = []
for _, row in group.iterrows():
    class_id = class_map[row['class_name']]
    x_center = (row['x_min'] + row['x_max']) / 2 / width
    y_center = (row['y_min'] + row['y_max']) / 2 / height
    bbox_width = (row['x_max'] - row['x_min']) / width
    bbox_height = (row['y_max'] - row['y_min']) / height
    yolo_lines.append(f"{class_id} {x_center} {y_center} {bbox_width} {bbox_height}")

# Save label file
label_path = os.path.join(yolo_labels_dir, split, file_name.replace(".jpg", ".txt"))
with open(label_path, "w") as f:
    f.write("\n".join(yolo_lines))
```

	image_id	class_name	x_min	y_min	x_max	y_max
0	0	pickup_truck	213	34	255	50
1	0	car	194	78	273	122
2	0	car	155	27	183	35
3	0	articulated_truck	43	25	109	55
4	0	car	106	32	124	45
...	...	...	...	...	...	...
351544	110590	car	18	57	97	98
351545	110591	articulated_truck	2	71	690	351
351546	110592	pickup_truck	3	240	214	378
351547	110592	car	465	111	507	135
351548	110592	non-motorized_vehicle	197	187	318	269

[293274 rows x 6 columns]

```
In [ ]: from glob import glob
print("Train images:", len(glob(os.path.join(yolo_images_dir, 'train', '*.jpg'))))
print("Val images:", len(glob(os.path.join(yolo_images_dir, 'val', '*.jpg'))))
```

Train images: 4500  
Val images: 1126

```
In [ ]: #create the yolo config file
yolo_dataset_path = "./sample_data/yolo_dataset/"
yaml_path = "./sample_data/yolo_dataset/data.yaml"
with open(yaml_path, "w") as f:
    f.write(f"""
path: {yolo_dataset_path}
train: images/train
val: images/val

nc: {len(class_map)}
names: {list(class_map.keys())}
""")
```

3. Create a CNN architecture for object detection of your choice to train an object detection model. Please note that algorithm or architecture selection is a very important aspect of ML model training, and you must pick the one that works the best for your dataset

```
In [ ]: model = YOLO("yolov8n.pt") # Use yolov8n for speed, or yolov8s/m/l/x for accuracy
results = model.train(data=yaml_path, epochs=20, imgsz=640)
```

Ultralytics 8.3.141 Python-3.11.12 torch-2.6.0+cu124 CUDA:0 (Tesla T4, 15095MiB)  
**engine/trainer:** agnostic\_nms=False, amp=True, augment=False, auto\_augment=randaugument, batch=16, bgr=0.0, box=7.5, cache=False, cfg=None, classes=None, close\_mosaic=10, cls=0.5, conf=None, copy\_paste=0.0, copy\_paste\_mode=flip, cos\_lr=False, cutmix=0.0, data=./sample\_data/yolo\_dataset/data.yaml, degrees=0.0, deterministic=True, device=None, dfl=1.5, dnn=False, dropout=0.0, dynamic=False, embed=None, epochs=20, erasing=0.4, exist\_ok=False, fliplr=0.5, flipud=0.0, format=torchscript, fraction=1.0, freeze=None, half=False, hsv\_h=0.015, hsv\_s=0.7, hsv\_v=0.4, imgsz=640, int8=False, iou=0.7, keras=False, kobj=1.0, line\_width=None, lr0=0.01, lrf=0.01, mask\_ratio=4, max\_det=300, mixup=0.0, mode=train, model=yolov8n.pt, momentum=0.937, mosaic=1.0, multi\_scale=False, name=train2, nbs=64, nms=False, opset=None, optimize=False, optimizer=auto, overlap\_mask=True, patience=100, perspective=0.0, plots=True, pose=12.0, pretrained=True, profile=False, project=None, rect=False, resume=False, retina\_masks=False, save=True, save\_conf=False, save\_crop=False, save\_dir=runs/detect/train2, save\_frames=False, save\_json=False, save\_period=-1, save\_txt=False, scale=0.5, seed=0, shear=0.0, show=False, show\_boxes=True, show\_conf=True, show\_labels=True, simplify=True, single\_cls=False, source=None, split=val, stream\_buffer=False, task=detect, time=None, tracker=botsort.yaml, translate=0.1, val=True, verbose=True, vid\_stride=1, visualize=False, warmup\_bias\_lr=0.1, warmup\_epochs=3.0, warmup\_momentum=0.8, weight\_decay=0.0005, workers=8, workspace=None  
Overriding model.yaml nc=80 with nc=11

		from	n	params	module	arguments
0		-1	1	464	ultralytics.nn.modules.conv.Conv	[3, 16, 3, 2]
1		-1	1	4672	ultralytics.nn.modules.conv.Conv	[16, 32, 3, 2]
2		-1	1	7360	ultralytics.nn.modules.block.C2f	[32, 32, 1, True]
3		-1	1	18560	ultralytics.nn.modules.conv.Conv	[32, 64, 3, 2]
4		-1	2	49664	ultralytics.nn.modules.block.C2f	[64, 64, 2, True]
5		-1	1	73984	ultralytics.nn.modules.conv.Conv	[64, 128, 3, 2]
6		-1	2	197632	ultralytics.nn.modules.block.C2f	[128, 128, 2, True]
7		-1	1	295424	ultralytics.nn.modules.conv.Conv	[128, 256, 3, 2]
8		-1	1	460288	ultralytics.nn.modules.block.C2f	[256, 256, 1, True]
9		-1	1	164608	ultralytics.nn.modules.block.SPPF	[256, 256, 5]
10		-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
11	[-1, 6]	-1	1	0	ultralytics.nn.modules.conv.Concat	[1]
12		-1	1	148224	ultralytics.nn.modules.block.C2f	[384, 128, 1]
13		-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
14	[-1, 4]	-1	1	0	ultralytics.nn.modules.conv.Concat	[1]
15		-1	1	37248	ultralytics.nn.modules.block.C2f	[192, 64, 1]
16		-1	1	36992	ultralytics.nn.modules.conv.Conv	[64, 64, 3, 2]
17	[-1, 12]	-1	1	0	ultralytics.nn.modules.conv.Concat	[1]
18		-1	1	123648	ultralytics.nn.modules.block.C2f	[192, 128, 1]
19		-1	1	147712	ultralytics.nn.modules.conv.Conv	[128, 128, 3, 2]
20	[-1, 9]	-1	1	0	ultralytics.nn.modules.conv.Concat	[1]
21		-1	1	493056	ultralytics.nn.modules.block.C2f	[384, 256, 1]
22	[15, 18, 21]	-1	1	753457	ultralytics.nn.modules.head.Detect	[11, [64, 128, 256]]

Model summary: 129 layers, 3,012,993 parameters, 3,012,977 gradients, 8.2 GFLOPs

Transferred 319/355 items from pretrained weights

Freezing layer 'model.22.dfl.conv.weight'

**AMP:** running Automatic Mixed Precision (AMP) checks...

**AMP:** checks passed ✓

**train:** Fast image access ✓ (ping: 0.0±0.0 ms, read: 815.5±333.6 MB/s, size: 24.9 KB)

**train:** Scanning /content/sample\_data/yolo\_dataset/labels/train.cache... 4500 images, 0 backgrounds, 0 corrupt: 100%|██████████| 4500/4500 [00:00<?, ?it/s]

**augmentations:** Blur(p=0.01, blur\_limit=(3, 7)), MedianBlur(p=0.01, blur\_limit=(3, 7)), ToGray(p=0.01, method='weighted\_average', num\_output\_channels=3), CLAHE(p=0.01, clip\_limit=(1.0, 4.0), tile\_grid\_size=(8, 8))

**val:** Fast image access ✓ (ping: 0.0±0.0 ms, read: 634.1±367.1 MB/s, size: 23.5 KB)

**val:** Scanning /content/sample\_data/yolo\_dataset/labels/val.cache... 1126 images, 0 backgrounds, 0 corrupt: 100%|██████████| 1126/1126 [00:00<?, ?it/s]

Plotting labels to runs/detect/train2/labels.jpg...

**optimizer:** 'optimizer=auto' found, ignoring 'lr0=0.01' and 'momentum=0.937' and determining best 'optimizer', 'lr0' and 'momentum' automatically...

**optimizer:** AdamW(lr=0.000667, momentum=0.9) with parameter groups 57 weight(decay=0.0), 64 weight(decay=0.0005), 63 bias(decay=0.0)

Image sizes 640 train, 640 val

Using 2 dataloader workers

Logging results to **runs/detect/train2**

Starting training for 20 epochs...

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
1/20	2.16G	1.231	2.372	1.08	14	640: 100% ██████████  282/282 [01:22<00:00, 3.43it/s]
		Class	Images	Instances	Box(P	R
		[00:11<00:00, 3.27it/s]				mAP50 mAP50-95): 100% ██████████  36/36
		all	1126	3806	0.451	0.228
					0.215	0.155
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
2/20	2.3G	1.16	1.595	1.044	23	640: 100% ██████████  282/282 [01:33<00:00, 3.00it/s]
		Class	Images	Instances	Box(P	R
		[00:09<00:00, 3.98it/s]				mAP50 mAP50-95): 100% ██████████  36/36
		all	1126	3806	0.394	0.302
					0.285	0.194
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size

3/20 0:00, 3.58it/s]	2.3G	1.14	1.403	1.04	25	640: 100% ██████████  282/282 [01:18<0		
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  36/36		
[00:08<00:00, 4.02it/s]	all	1126	3806	0.462	0.365	0.343	0.232	
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size		
4/20 0:00, 3.62it/s]	2.3G	1.115	1.259	1.032	22	640: 100% ██████████  282/282 [01:17<0		
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  36/36		
[00:08<00:00, 4.32it/s]	all	1126	3806	0.541	0.417	0.419	0.291	
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size		
5/20 0:00, 3.60it/s]	2.3G	1.08	1.151	1.019	16	640: 100% ██████████  282/282 [01:18<0		
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  36/36		
[00:09<00:00, 3.89it/s]	all	1126	3806	0.439	0.437	0.41	0.279	
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size		
6/20 0:00, 3.69it/s]	2.3G	1.068	1.086	1.013	20	640: 100% ██████████  282/282 [01:16<0		
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  36/36		
[00:09<00:00, 3.80it/s]	all	1126	3806	0.418	0.526	0.452	0.319	
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size		
7/20 0:00, 3.73it/s]	2.31G	1.037	1.018	1.003	28	640: 100% ██████████  282/282 [01:15<0		
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  36/36		
[00:08<00:00, 4.38it/s]	all	1126	3806	0.55	0.463	0.475	0.334	
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size		
8/20 0:00, 3.67it/s]	2.33G	1.04	0.9863	1.004	22	640: 100% ██████████  282/282 [01:16<0		
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  36/36		
[00:08<00:00, 4.43it/s]	all	1126	3806	0.554	0.433	0.46	0.318	
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size		
9/20 0:00, 3.69it/s]	2.33G	1.024	0.9512	0.9983	17	640: 100% ██████████  282/282 [01:16<0		
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  36/36		
[00:09<00:00, 3.90it/s]	all	1126	3806	0.478	0.538	0.51	0.357	
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size		
10/20 0:00, 3.71it/s]	2.33G	1.014	0.9296	0.9955	26	640: 100% ██████████  282/282 [01:16<0		
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  36/36		
[00:08<00:00, 4.08it/s]	all	1126	3806	0.474	0.55	0.528	0.378	
Closing dataloader mosaic								
alumentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limit=(3, 7)), ToGray(p=0.01, method='weighted_average', num_output_channels=3), CLAHE(p=0.01, clip_limit=(1.0, 4.0), tile_grid_size=(8, 8))								
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size		
11/20 0:00, 3.74it/s]	2.34G	0.9834	0.8624	0.9731	8	640: 100% ██████████  282/282 [01:15<0		
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  36/36		
[00:08<00:00, 4.31it/s]	all	1126	3806	0.571	0.517	0.554	0.392	
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size		
12/20 0:00, 3.88it/s]	2.37G	0.9693	0.8148	0.9642	11	640: 100% ██████████  282/282 [01:12<0		
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  36/36		
[00:08<00:00, 4.05it/s]	all	1126	3806	0.581	0.554	0.562	0.396	
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size		
13/20 0:00, 3.92it/s]	2.37G	0.9593	0.7927	0.9597	18	640: 100% ██████████  282/282 [01:12<0		
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  36/36		
[00:08<00:00, 4.12it/s]	all	1126	3806	0.507	0.582	0.544	0.389	
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size		

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
14/20 0:00, 3.87it/s]	2.37G	0.9468	0.767	0.9581	10	640: 100% ██████████  282/282 [01:12<00:00, 3.87it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  36/36
[00:08<00:00, 4.17it/s]	all	1126	3806	0.524	0.575	0.568 0.408
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
15/20 0:00, 3.84it/s]	2.37G	0.9396	0.7394	0.9553	11	640: 100% ██████████  282/282 [01:13<00:00, 3.84it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  36/36
[00:09<00:00, 3.95it/s]	all	1126	3806	0.568	0.571	0.571 0.413
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
16/20 0:00, 3.87it/s]	2.37G	0.9285	0.7166	0.9487	13	640: 100% ██████████  282/282 [01:12<00:00, 3.87it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  36/36
[00:08<00:00, 4.42it/s]	all	1126	3806	0.574	0.598	0.597 0.426
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
17/20 0:00, 3.89it/s]	2.37G	0.9121	0.6981	0.945	9	640: 100% ██████████  282/282 [01:12<00:00, 3.89it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  36/36
[00:08<00:00, 4.02it/s]	all	1126	3806	0.551	0.617	0.588 0.428
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
18/20 0:00, 3.93it/s]	2.37G	0.9094	0.6752	0.9408	12	640: 100% ██████████  282/282 [01:11<00:00, 3.93it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  36/36
[00:08<00:00, 4.06it/s]	all	1126	3806	0.618	0.581	0.601 0.441
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
19/20 0:00, 3.89it/s]	2.37G	0.8942	0.6542	0.9312	7	640: 100% ██████████  282/282 [01:12<00:00, 3.89it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  36/36
[00:08<00:00, 4.50it/s]	all	1126	3806	0.608	0.619	0.615 0.448
Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances	Size
20/20 0:00, 3.85it/s]	2.37G	0.8829	0.6359	0.9299	7	640: 100% ██████████  282/282 [01:13<00:00, 3.85it/s]
Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  36/36
[00:08<00:00, 4.13it/s]	all	1126	3806	0.625	0.59	0.616 0.452

20 epochs completed in 0.474 hours.  
Optimizer stripped from runs/detect/train2/weights/last.pt, 6.2MB  
Optimizer stripped from runs/detect/train2/weights/best.pt, 6.2MB

Validating runs/detect/train2/weights/best.pt...  
Ultralytics 8.3.141 Python-3.11.12 torch-2.6.0+cu124 CUDA:0 (Tesla T4, 15095MiB)  
Model summary (fused): 72 layers, 3,007,793 parameters, 0 gradients, 8.1 GFLOPs

Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████  36/36
[00:10<00:00, 3.49it/s]	all	1126	3806	0.624	0.589	0.615 0.451
pickup_truck	311	434	0.795	0.786	0.844	0.688
car	944	2544	0.835	0.882	0.912	0.667
articulated_truck	70	81	0.655	0.751	0.725	0.542
bus	105	121	0.931	0.897	0.938	0.83
motorized_vehicle	232	303	0.57	0.363	0.408	0.245
work_van	99	104	0.629	0.654	0.615	0.5
single_unit_truck	57	60	0.558	0.55	0.518	0.357
pedestrian	45	92	0.457	0.266	0.331	0.166
bicycle	22	26	0.481	0.577	0.64	0.408
non-motorized_vehicle	19	19	0.218	0.105	0.0913	0.0615
motorcycle	20	22	0.739	0.645	0.742	0.497

Speed: 0.2ms preprocess, 1.7ms inference, 0.0ms loss, 1.8ms postprocess per image  
Results saved to **runs/detect/train2**

#### 4. Evaluate the model and check the test results

```
In [ ]: import matplotlib.pyplot as plt
import pandas as pd

# Replace with the actual path to your results.csv
results_path = './runs/detect/train/results.csv'

# Read the CSV, skipping initial comments (if any)
# You might need to adjust 'skipinitialspace' or 'header' based on your file's exact format
try:
```

```


df = pd.read_csv(results_path)
except FileNotFoundError:
    print(f"Error: {results_path} not found. Please check the path.")
    exit()


# Clean column names (YOLOv8 often has leading spaces)
df.columns = [col.strip() for col in df.columns]

# Plotting Loss
plt.figure(figsize=(12, 6))
plt.plot(df['epoch'], df['train/box_loss'] + df['train/cls_loss'] + df['train/dfl_loss'], label='Train Loss (Combined)')
plt.plot(df['epoch'], df['val/box_loss'] + df['val/cls_loss'] + df['val/dfl_loss'], label='Val Loss (Combined)')
plt.title('Training and Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.show()

# Plotting mAP
plt.figure(figsize=(12, 6))
plt.plot(df['epoch'], df['metrics/mAP50(B)'], label='mAP50(B) - Validation')
plt.plot(df['epoch'], df['metrics/mAP50-95(B)'], label='mAP50-95(B) - Validation')
plt.title('Validation mAP Metrics')
plt.xlabel('Epoch')
plt.ylabel('mAP')
plt.legend()
plt.grid(True)
plt.show()

```





5. Run inferences on sample images and see if vehicles are detected accurately

```
In [ ]: results_in = model.predict("./sample_data/yolo_dataset/images/val/00004523.jpg", save=True, conf=0.25)
print(results_in)

# Extract first result
r = results_in[0]

# Show class names
print("Class Names:", model.names)

# Print all predictions
for box, conf, cls_id in zip(r.bboxes.xyxy, r.bboxes.conf, r.bboxes.cls):
    print(f"Class: {model.names[int(cls_id)]}, Confidence: {conf:.2f}, Box: {box.tolist()}")
```

```

image 1/1 /content/sample_data/yolo_dataset/images/val/00004523.jpg: 448x640 5 cars, 1 work_van, 121.2ms
Speed: 3.0ms preprocess, 121.2ms inference, 2.2ms postprocess per image at shape (1, 3, 448, 640)
Results saved to runs/detect/train22
[ultralytics.engine.results.Results object with attributes:

boxes: ultralytics.engine.results.Boxes object
keypoints: None
masks: None
names: {0: 'pickup_truck', 1: 'car', 2: 'articulated_truck', 3: 'bus', 4: 'motorized_vehicle', 5: 'work_van', 6:
'single_unit_truck', 7: 'pedestrian', 8: 'bicycle', 9: 'non-motorized_vehicle', 10: 'motorcycle'}
obb: None
orig_img: array([[[ 31, 12, 9],
[ 37, 18, 13],
[ 24, 6, 0],
...,
[ 31, 14, 17],
[ 22, 8, 10],
[ 24, 10, 12]],
[[ 18, 0, 0],
[ 66, 47, 42],
[ 95, 77, 70],
...,
[ 36, 19, 22],
[ 19, 5, 7],
[ 21, 7, 9]],
[[ 58, 39, 36],
[166, 147, 142],
[240, 222, 215],
...,
[ 30, 16, 18],
[ 13, 0, 1],
[ 19, 5, 7]],
...,
[[ 1, 17, 16],
[ 26, 42, 41],
[ 50, 65, 67],
...,
[ 0, 0, 0],
[ 0, 0, 0],
[ 0, 0, 0]],
[[ 3, 19, 18],
[ 27, 43, 42],
[ 49, 64, 66],
...,
[ 0, 0, 0],
[ 0, 0, 0],
[ 0, 0, 0]],
[[ 3, 19, 18],
[ 28, 44, 43],
[ 49, 64, 66],
...,
[ 0, 0, 0],
[ 0, 0, 0],
[ 0, 0, 0]]], dtype=uint8)
orig_shape: (228, 342)
path: '/content/sample_data/yolo_dataset/images/val/00004523.jpg'
probs: None
save_dir: 'runs/detect/train22'
speed: {'preprocess': 2.9612039998028195, 'inference': 121.22001700026885, 'postprocess': 2.164272000300116}]
Class Names: {0: 'pickup_truck', 1: 'car', 2: 'articulated_truck', 3: 'bus', 4: 'motorized_vehicle', 5: 'work_van', 6: 'single_unit_truck', 7: 'pedestrian', 8: 'bicycle', 9: 'non-motorized_vehicle', 10: 'motorcycle'}
Class: work_van, Confidence: 0.80, Box: [69.21305847167969, 55.57606506347656, 138.12481689453125, 89.46800994873047]
Class: car, Confidence: 0.75, Box: [151.51165771484375, 27.57723617553711, 182.0041046142578, 35.35712432861328]
Class: car, Confidence: 0.63, Box: [69.44784545898438, 55.727657318115234, 138.2427215576172, 89.75243377685547]
Class: car, Confidence: 0.61, Box: [226.2789764404297, 27.015281677246094, 247.7737274169922, 35.183013916015625]
Class: car, Confidence: 0.58, Box: [110.9125747680664, 24.340269088745117, 129.16390991210938, 31.648591995239258]
Class: car, Confidence: 0.34, Box: [27.07554054260254, 43.211185455322266, 45.96173095703125, 55.497798919677734]
]

```

```

In [ ]: # r.plot() returns an annotated image (numpy array)
annotated_img = r.plot()

# Display it with matplotlib
plt.figure(figsize=(10, 8))

```



```
plt.imshow(annotated_img)
plt.axis('off')
plt.title("YOLOv8 Prediction")
plt.show()
```

YOLOv8 Prediction



## Part 2

### Data Science

#### 1. Preliminary data inspection and cleaning

- Perform preliminary data inspection, checking for data types, missing values, and duplicates
- Remove any columns that might not be relevant for the analysis

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [ ]: # Load the dataset
df_full = pd.read_csv("./sample_data/Tesla - Deaths.csv")

df = df_full.iloc[:294, :]
col_maxima = df.tail(10)
print(col_maxima)
#shape
print("Dataset Shape:", df.shape)
# Display basic info
print("Dataset Info:")
print(df.info())

# Show first few rows
print("\nSample Data:")
print(df.head())

# Check for missing values
print("\nMissing Values:")
print(df.isnull().sum())

# Check for duplicates
print("\nDuplicate Rows:")
print(df.duplicated().sum())

df.columns = df.columns.str.strip()
print("Cleaned column names:")
print(df.columns.tolist())
```

```

# Define columns to drop (you can customize this)
irrelevant_cols = [
    "Note", "Source", "Unnamed: 16", "Unnamed: 17"
]

# Drop the columns if they exist in the dataset
df = df.drop(columns=[col for col in irrelevant_cols if col in df.columns])

# Verify cleanup
print("\nCleaned DataFrame:")
print(df.columns)

# Identify numeric and object columns
num_cols = df.select_dtypes(include=['int64', 'float64']).columns
str_cols = df.select_dtypes(include=['object']).columns
print(num_cols)
#fill empty columns with NaN
df[df.columns] = df[df.columns].replace(['-', ' - ', ''], np.nan)

df.tail(100)

```

	Case #	Year	Date	Country	State	\
284	10.0	2015.0	12/22/2015	Canada	-	
285	9.0	2015.0	11/18/2015	USA	CA	
286	8.0	2015.0	6/22/2015	USA	CA	
287	7.0	2015.0	1/22/2015	USA	CA	
288	6.0	2014.0	12/30/2014	USA	CA	
289	5.0	2014.0	7/14/2014	USA	CA	
290	4.0	2014.0	7/4/2014	USA	CA	
291	3.0	2014.0	7/4/2014	USA	CA	
292	2.0	2013.0	11/2/2013	USA	CA	
293	1.0	2013.0	4/2/2013	USA	CA	

	Description	Deaths	Tesla driver	\
284	Struck by dumptruck	1.0	1	
285	Tesla kills pedestrian	1.0	-	
286	Tesla drives off cliff	1.0	1	
287	Tesla drives off cliff	1.0	1	
288	Tesla drives off cliff	1.0	1	
289	Tesla kills motorcyclist	1.0	-	
290	Thief crashes stolen Tesla	1.0	1	
291	Tesla rear ends stopped car	3.0	-	
292	Tesla kills cyclist	1.0	-	
293	Tesla veers into opposite lane	2.0	-	

	Tesla occupant	Other vehicle	...	Verified Tesla Autopilot Deaths	\
284	-	-	...	-	
285	-	-	...	-	
286	-	-	...	-	
287	-	-	...	-	
288	-	-	...	-	
289	-	1	...	-	
290	-	-	...	-	
291	-	3	...	-	
292	-	-	...	-	
293	-	2	...	-	

	Verified Tesla Autopilot Deaths + All Deaths Reported to NHTSA SGO	\
284	-	
285	-	
286	-	
287	-	
288	-	
289	-	
290	-	
291	-	
292	-	
293	-	

	Unnamed: 16	\
284	<a href="https://web.archive.org/web/20220817120756/ht...">https://web.archive.org/web/20220817120756/ht...</a>	
285	<a href="https://web.archive.org/web/20220817120754/ht...">https://web.archive.org/web/20220817120754/ht...</a>	
286	<a href="https://web.archive.org/web/20220817120755/ht...">https://web.archive.org/web/20220817120755/ht...</a>	
287	<a href="https://web.archive.org/web/20220817120837/ht...">https://web.archive.org/web/20220817120837/ht...</a>	
288	<a href="https://web.archive.org/web/20220817120806/ht...">https://web.archive.org/web/20220817120806/ht...</a>	
289	<a href="https://web.archive.org/web/20220817120807/ht...">https://web.archive.org/web/20220817120807/ht...</a>	
290	<a href="https://web.archive.org/web/20220817120839/ht...">https://web.archive.org/web/20220817120839/ht...</a>	
291	<a href="https://web.archive.org/web/20220412004559/ht...">https://web.archive.org/web/20220412004559/ht...</a>	
292	<a href="https://web.archive.org/web/20220817121049/ht...">https://web.archive.org/web/20220817121049/ht...</a>	
293	<a href="https://web.archive.org/web/20150425055520/ht...">https://web.archive.org/web/20150425055520/ht...</a>	

Unnamed: 17 \

```

284 https://web.archive.org/web/20220817120756/ht...
285 https://web.archive.org/web/20220817120754/ht...
286 https://web.archive.org/web/20220817120755/ht...
287 https://web.archive.org/web/20220817120837/ht...
288 https://web.archive.org/web/20220817120806/ht...
289 https://web.archive.org/web/20220817120807/ht...
290 https://web.archive.org/web/20220817120839/ht...
291 https://web.archive.org/web/20220412004559/ht...
292 https://web.archive.org/web/20220817121049/ht...
293 https://web.archive.org/web/20150425055520/ht...

```

	Source	Note	\
284	https://web.archive.org/web/20220817120756/ht...	NaN	
285	https://web.archive.org/web/20220817120754/ht...	NaN	
286	https://web.archive.org/web/20220817120755/ht...	NaN	
287	https://web.archive.org/web/20220817120837/ht...	NaN	
288	https://web.archive.org/web/20220817120806/ht...	NaN	
289	https://web.archive.org/web/20220817120807/ht...	NaN	
290	https://web.archive.org/web/20220817120839/ht...	NaN	
291	https://web.archive.org/web/20220412004559/ht...	NaN	
292	https://web.archive.org/web/20220817121049/ht...	NaN	
293	https://web.archive.org/web/20150425055520/ht...	NaN	

	Deceased 1	Deceased 2	Deceased 3	\
284	NaN	NaN	NaN	
285	NaN	NaN	NaN	
286	Tim Devine	NaN	NaN	
287	Peter Kleis	NaN	NaN	
288	Louis Francis Thoeleck	NaN	NaN	
289	NaN	NaN	NaN	
290	Joshua Slot	NaN	NaN	
291	NaN	NaN	NaN	
292	NaN	NaN	NaN	
293	Alberto Casique-Salinas	Armando Garcia-Gonzales	NaN	

	Deceased 4
284	NaN
285	NaN
286	NaN
287	NaN
288	NaN
289	NaN
290	NaN
291	NaN
292	NaN
293	NaN

```

[10 rows x 24 columns]
Dataset Shape: (294, 24)
Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 294 entries, 0 to 293
Data columns (total 24 columns):

```

#	Column	Non-Null Count	Dtype
0	Case #	294 non-null	float64
1	Year	294 non-null	float64
2	Date	294 non-null	object
3	Country	294 non-null	object
4	State	294 non-null	object
5	Description	294 non-null	object
6	Deaths	294 non-null	float64
7	Tesla driver	289 non-null	object
8	Tesla occupant	285 non-null	object
9	Other vehicle	290 non-null	object
10	Cyclists/ Peds	291 non-null	object
11	TSLA+cycl / peds	292 non-null	object
12	Model	294 non-null	object
13	Autopilot claimed	276 non-null	object
14	Verified Tesla Autopilot Deaths	290 non-null	object
15	Verified Tesla Autopilot Deaths + All Deaths Reported to NHTSA SGO	293 non-null	object
16	Unnamed: 16	289 non-null	object
17	Unnamed: 17	289 non-null	object
18	Source	294 non-null	object
19	Note	9 non-null	object
20	Deceased 1	87 non-null	object
21	Deceased 2	17 non-null	object
22	Deceased 3	4 non-null	object
23	Deceased 4	0 non-null	float64

```

dtypes: float64(4), object(20)
memory usage: 55.3+ KB
None

```

Sample Data:

Case #	Year	Date	Country	State	\
0	294.0	2022.0	1/17/2023	USA	CA
1	293.0	2022.0	1/7/2023	Canada	-
2	292.0	2022.0	1/7/2023	USA	WA
3	291.0	2022.0	12/22/2022	USA	GA
4	290.0	2022.0	12/19/2022	Canada	-

	Description	Deaths	Tesla driver	\
0	Tesla crashes into back of semi	1.0	1	
1	Tesla crashes	1.0	1	
2	Tesla hits pole, catches on fire	1.0	-	
3	Tesla crashes and burns	1.0	1	
4	Tesla crashes into storefront	1.0	-	

	Tesla occupant	Other vehicle	...	Verified Tesla Autopilot Deaths	\
0	-	-	...	-	
1	-	-	...	-	
2	1	-	...	-	
3	-	-	...	-	
4	-	-	...	-	

	Verified Tesla Autopilot Deaths + All Deaths Reported to NHTSA SGO	\
0	-	
1	-	
2	-	
3	-	
4	-	

	Unnamed: 16	\
0	<a href="https://web.archive.org/web/20221222203930/ht...">https://web.archive.org/web/20221222203930/ht...</a>	
1	<a href="https://web.archive.org/web/20221222203930/ht...">https://web.archive.org/web/20221222203930/ht...</a>	
2	<a href="https://web.archive.org/web/20221222203930/ht...">https://web.archive.org/web/20221222203930/ht...</a>	
3	<a href="https://web.archive.org/web/20221222203930/ht...">https://web.archive.org/web/20221222203930/ht...</a>	
4	<a href="https://web.archive.org/web/20221223203725/ht...">https://web.archive.org/web/20221223203725/ht...</a>	

	Unnamed: 17	\
0	<a href="https://web.archive.org/web/20221222203930/ht...">https://web.archive.org/web/20221222203930/ht...</a>	
1	<a href="https://web.archive.org/web/20221222203930/ht...">https://web.archive.org/web/20221222203930/ht...</a>	
2	<a href="https://web.archive.org/web/20221222203930/ht...">https://web.archive.org/web/20221222203930/ht...</a>	
3	<a href="https://web.archive.org/web/20221222203930/ht...">https://web.archive.org/web/20221222203930/ht...</a>	
4	<a href="https://web.archive.org/web/20221223203725/ht...">https://web.archive.org/web/20221223203725/ht...</a>	

	Source	Note	\
0	<a href="https://web.archive.org/web/20230118162813/ht...">https://web.archive.org/web/20230118162813/ht...</a>	NaN	
1	<a href="https://web.archive.org/web/20230109041434/ht...">https://web.archive.org/web/20230109041434/ht...</a>	NaN	
2	<a href="https://web.archive.org/web/20230107232745/ht...">https://web.archive.org/web/20230107232745/ht...</a>	NaN	
3	<a href="https://web.archive.org/web/20221222203930/ht...">https://web.archive.org/web/20221222203930/ht...</a>	NaN	
4	<a href="https://web.archive.org/web/20221223203725/ht...">https://web.archive.org/web/20221223203725/ht...</a>	NaN	

	Deceased 1	Deceased 2	Deceased 3	Deceased 4
0	NaN	NaN	NaN	NaN
1	Taren Singh Lal	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN

[5 rows x 24 columns]

Missing Values:

Case #	0
Year	0
Date	0
Country	0
State	0
Description	0
Deaths	0
Tesla driver	5
Tesla occupant	9
Other vehicle	4
Cyclists/ Peds	3
TSLA+cycl / peds	2
Model	0
Autopilot claimed	18
Verified Tesla Autopilot Deaths	4
Verified Tesla Autopilot Deaths + All Deaths Reported to NHTSA SGO	1
Unnamed: 16	5
Unnamed: 17	5
Source	0
Note	285
Deceased 1	207
Deceased 2	277
Deceased 3	290

Deceased 4  
dtype: int64

294

Duplicate Rows:

0

Cleaned column names:

```
['Case #', 'Year', 'Date', 'Country', 'State', 'Description', 'Deaths', 'Tesla driver', 'Tesla occupant', 'Other vehicle', 'Cyclists/ Peds', 'TSLA+cycl / peds', 'Model', 'Autopilot claimed', 'Verified Tesla Autopilot Deaths', 'Verified Tesla Autopilot Deaths + All Deaths Reported to NHTSA SGO', 'Unnamed: 16', 'Unnamed: 17', 'Source', 'Note', 'Deceased 1', 'Deceased 2', 'Deceased 3', 'Deceased 4']
```

Cleaned DataFrame:

```
Index(['Case #', 'Year', 'Date', 'Country', 'State', 'Description', 'Deaths', 'Tesla driver', 'Tesla occupant', 'Other vehicle', 'Cyclists/ Peds', 'TSLA+cycl / peds', 'Model', 'Autopilot claimed', 'Verified Tesla Autopilot Deaths', 'Verified Tesla Autopilot Deaths + All Deaths Reported to NHTSA SGO', 'Deceased 1', 'Deceased 2', 'Deceased 3', 'Deceased 4'], dtype='object')
```

```
Index(['Case #', 'Year', 'Deaths', 'Deceased 4'], dtype='object')
```

Out[ ]:

	Case #	Year	Date	Country	State	Description	Deaths	Tesla driver	Tesla occupant	Other vehicle	Cyclists/ Peds	TSLA+cycl / peds	Model	Autopilot claim
194	100.0	2019.0	12/29/2019	USA	CA	Tesla runs red light after exiting freeway	2.0	NaN	NaN	2	NaN	NaN	S	
195	99.0	2019.0	12/23/2019	Slovenia	NaN	Drunk driver strikes Tesla	1.0	NaN	NaN	1	NaN	NaN	S	N
196	98.0	2019.0	12/19/2019	USA	FL	Police SUV strikes Tesla	2.0	1	1	NaN	NaN	2	NaN	N
197	97.0	2019.0	12/11/2019	USA	CA	Tesla strikes turning Lexus	1.0	NaN	NaN	1	NaN	NaN	X	N
198	96.0	2019.0	12/10/2019	USA	NV	Kia rear ends Tesla	1.0	NaN	NaN	1	NaN	NaN	NaN	N
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
289	5.0	2014.0	7/14/2014	USA	CA	Tesla kills motorcyclist	1.0	NaN	NaN	1	NaN	NaN	NaN	N
290	4.0	2014.0	7/4/2014	USA	CA	Thief crashes stolen Tesla	1.0	1	NaN	NaN	NaN	1	NaN	N
291	3.0	2014.0	7/4/2014	USA	CA	Tesla rear ends stopped car	3.0	NaN	NaN	3	NaN	NaN	NaN	N
292	2.0	2013.0	11/2/2013	USA	CA	Tesla kills cyclist	1.0	NaN	NaN	NaN	1	1	NaN	N
293	1.0	2013.0	4/2/2013	USA	CA	Tesla veers into opposite lane	2.0	NaN	NaN	2	NaN	NaN	S	N

100 rows × 20 columns



## Exploratory Data Analysis

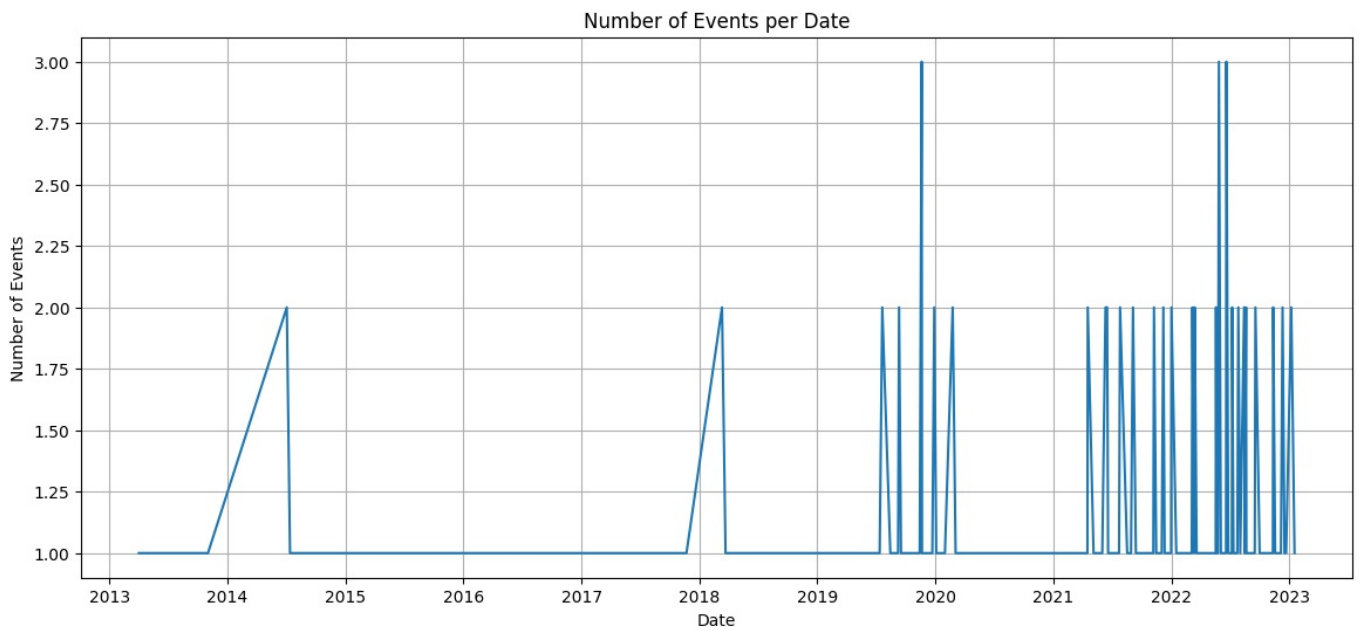
a. Perform an in-depth exploratory data analysis on the number of events by date, per year, and per day for each state and country

```
In [ ]: df['Date'] = pd.to_datetime(df['Date'], errors='coerce') # Convert Date to datetime, coerce errors
```

```
In [ ]: events_per_date = df.groupby('Date').size().reset_index(name='count')
```

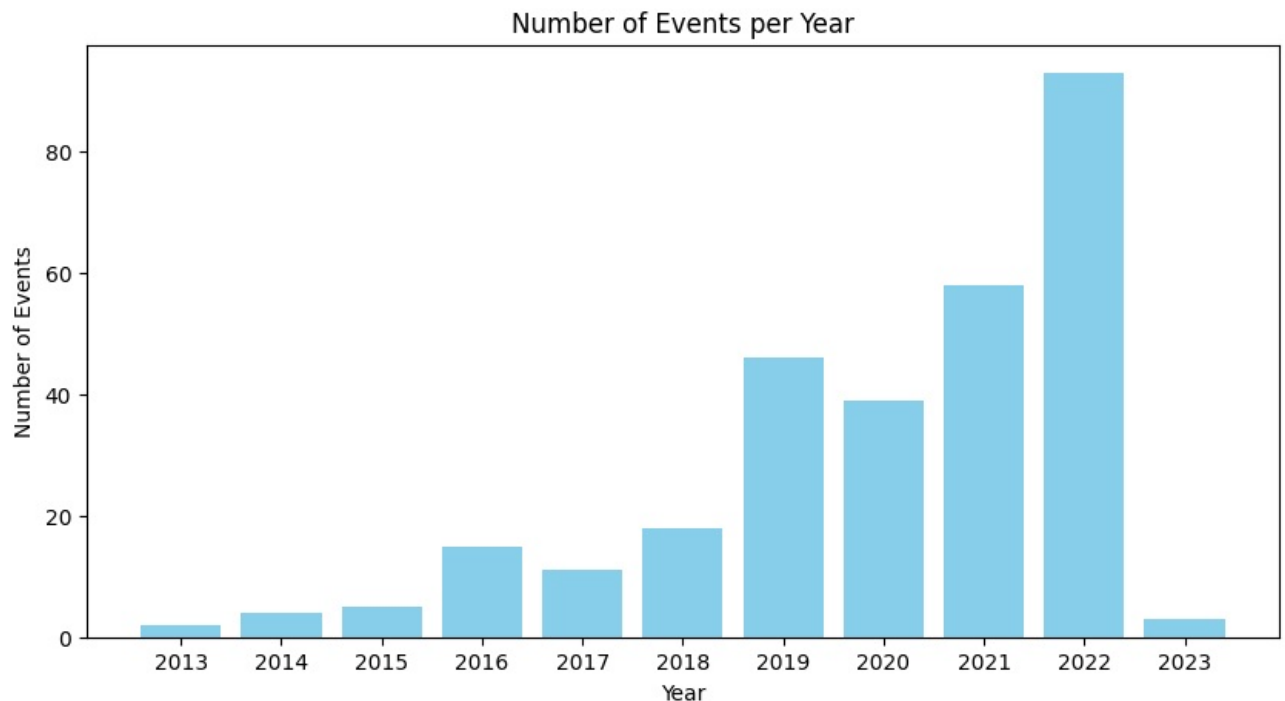
```
plt.figure(figsize=(14,6))
plt.plot(events_per_date['Date'], events_per_date['count'])
plt.title('Number of Events per Date')
```

```
plt.xlabel('Date')
plt.ylabel('Number of Events')
plt.grid(True)
plt.show()
```




```
In [63]: # Events count per year
df['Year'] = df['Date'].dt.year
events_per_year = df.groupby('Year').size().reset_index(name='count')

plt.figure(figsize=(10,5))
plt.bar(events_per_year['Year'], events_per_year['count'], color='skyblue')
plt.title('Number of Events per Year')
plt.xlabel('Year')
plt.ylabel('Number of Events')
plt.xticks(events_per_year['Year'])
plt.show()
```



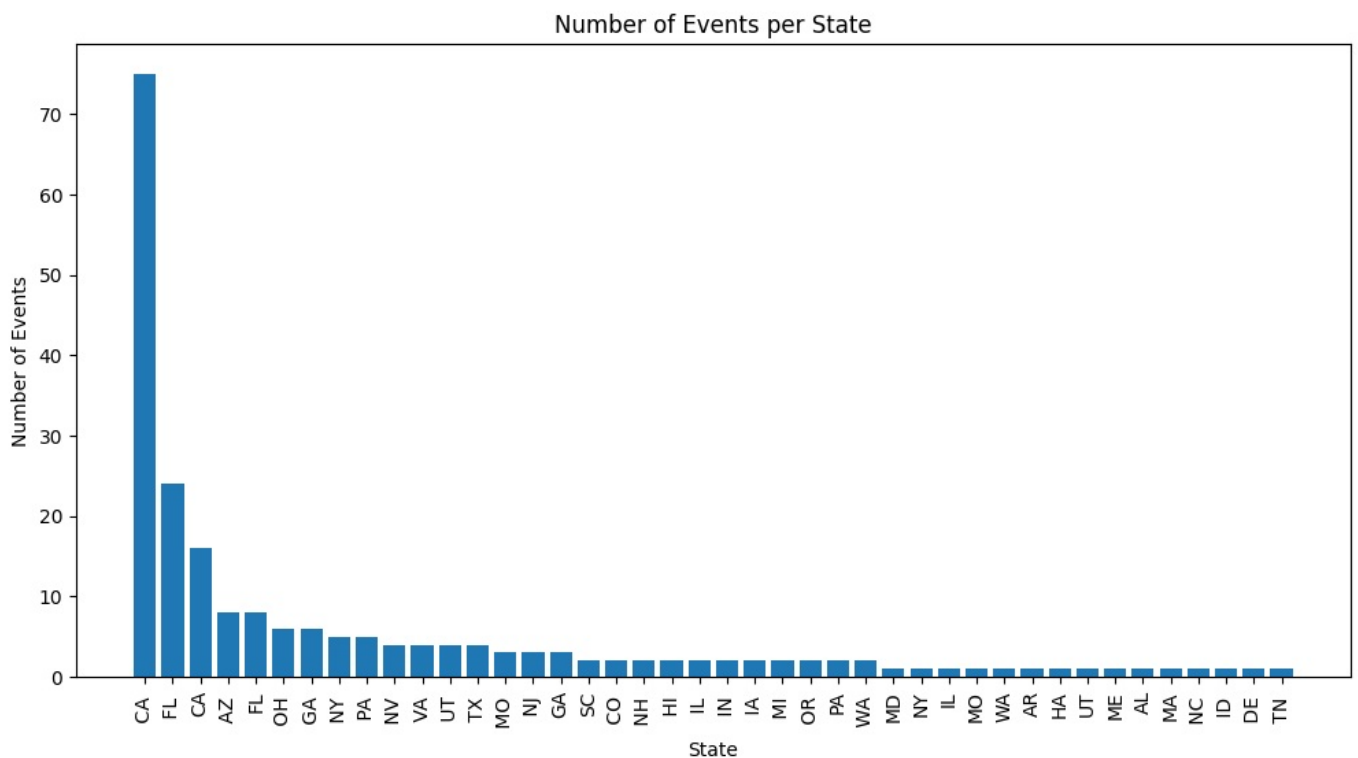
```
In [ ]: #Events count per day of week
df['DayOfWeek'] = df['Date'].dt.day_name()
events_per_day = df.groupby('DayOfWeek').size().reindex([
    'Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday'
]).reset_index(name='count')

plt.figure(figsize=(10,5))
plt.bar(events_per_day['DayOfWeek'], events_per_day['count'], color='coral')
plt.title('Number of Events per Day of Week')
plt.xlabel('Day of Week')
plt.ylabel('Number of Events')
plt.show()
```




```
In [ ]: #Events count per State
events_per_state = df['State'].value_counts().reset_index()
events_per_state.columns = ['State', 'count']

plt.figure(figsize=(12,6))
plt.bar(events_per_state['State'], events_per_state['count'])
plt.title('Number of Events per State')
plt.xlabel('State')
plt.ylabel('Number of Events')
plt.xticks(rotation=90)
plt.show()
```



```
In [ ]: #Events count per Country
events_per_country = df['Country'].value_counts().reset_index()
events_per_country.columns = ['Country', 'count']

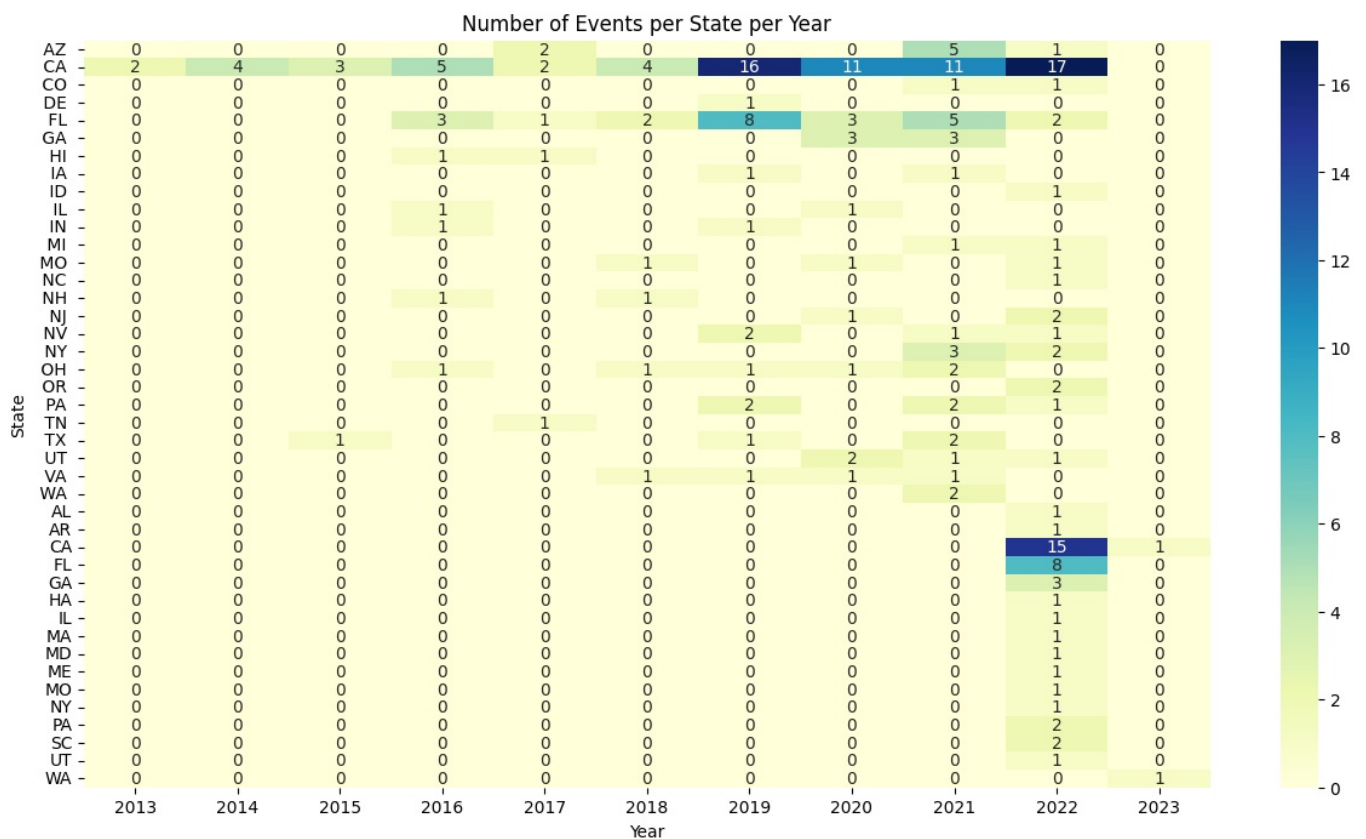
plt.figure(figsize=(8,5))
plt.bar(events_per_country['Country'], events_per_country['count'], color='green')
plt.title('Number of Events per Country')
plt.xlabel('Country')
plt.ylabel('Number of Events')
plt.xticks(rotation=45, ha='right')
plt.show()
```



In [ ]: #Heatmap of events per State and Year

```
state_year = df.groupby(['State', 'Year']).size().unstack(fill_value=0)

plt.figure(figsize=(15,8))
sns.heatmap(state_year, annot=True, fmt='d', cmap='YlGnBu')
plt.title('Number of Events per State per Year')
plt.xlabel('Year')
plt.ylabel('State')
plt.show()
```



b. Analyze the different aspects of the death events. For example:

- What is the number of victims (deaths) in each accident?



```
In [ ]: # Show accident ID and deaths
if 'Case #' in df.columns:
    print(df[['Case #', 'Deaths']])
else:
    print(df[['Deaths']])
```

	Case #	Deaths
0	294.0	1.0
1	293.0	1.0
2	292.0	1.0
3	291.0	1.0
4	290.0	1.0
...	...	...
289	5.0	1.0
290	4.0	1.0
291	3.0	3.0
292	2.0	1.0
293	1.0	2.0

[294 rows x 2 columns]

```
In [ ]: print("Deaths per accident statistics:")
print(df['Deaths'].describe())
df['Deaths'] = pd.to_numeric(df['Deaths'], errors='coerce')
# Count how many accidents with zero deaths vs non-zero deaths
print("\nAccidents with zero deaths:", (df['Deaths'] == 0).sum())
print("Accidents with one or more deaths:", (df['Deaths'] >= 1).sum())
```

Deaths per accident statistics:

```
count    294.000000
mean      1.200680
std       0.513171
min       1.000000
25%      1.000000
50%      1.000000
75%      1.000000
max       4.000000
Name: Deaths, dtype: float64
```

Accidents with zero deaths: 0

Accidents with one or more deaths: 294

- How many times did tesla drivers die ?

```
In [ ]: # Count number of deaths
tesla_driver_deaths = (df['Verified Tesla Autopilot Deaths'].dropna().astype(int) >= 1).sum()

print(f"Tesla drivers died in {tesla_driver_deaths} accident(s).")
```

Tesla drivers died in 16 accident(s).

- What is the proportion of events in which one or more occupants died ?

```
In [ ]: # Ensure 'Deaths' is numeric
df['Deaths'] = pd.to_numeric(df['Deaths'], errors='coerce')

# Total number of events
total_events = len(df)

# Events with at least one death
fatal_events = (df['Deaths'] >= 1).sum()

# Proportion calculation
proportion = fatal_events / total_events

print(f"Total Events: {total_events}")
print(f"Events with ≥1 Death: {fatal_events}")
print(f"Proportion: {proportion:.2%}")
```

Total Events: 294

Events with ≥1 Death: 294

Proportion: 100.00%

- What is the distribution of events in which the vehicle hit a cyclist or a pedestrian ?

```
In [ ]: df['Cyclists/ Peds'] = pd.to_numeric(df['Cyclists/ Peds'], errors='coerce')

print("Distribution of cyclist/pedestrian hits:")
print(df['Cyclists/ Peds'].value_counts().sort_index())
```

```
plt.figure(figsize=(8,5))
df['Cyclists/ Peds'].dropna().astype(int).value_counts().sort_index().plot(kind='bar', color='orange')
plt.title('Number of Events vs Cyclist/Pedestrian Hits')
plt.xlabel('Number of Cyclists/Pedestrians Hit')
plt.ylabel('Number of Events')
plt.xticks(rotation=0)
plt.grid(axis='y')
plt.tight_layout()
plt.show()
```


Distribution of cyclist/pedestrian hits:

Cyclists/ Peds

1.0 42

2.0 2

Name: count, dtype: int64



- How many times did the accident involve the death of an occupant or driver of a Tesla along with a cyclist or pedestrian?

```
In [ ]: # Convert relevant columns
print(df.columns)
df['Cyclists/ Peds'] = pd.to_numeric(df['Cyclists/ Peds'], errors='coerce')
df['Tesla driver'] = pd.to_numeric(df['Tesla driver'], errors='coerce')
df['Tesla occupant'] = pd.to_numeric(df['Tesla occupant'], errors='coerce')

tesla_death = (
    (df['Tesla driver'] >= 1) |
    (df['Tesla occupant'] >= 1)
)
cycl_peds_hit = df['Cyclists/ Peds'] >= 1
combined_condition = tesla_death & cycl_peds_hit
count = combined_condition.sum()

print(f"Number of accidents involving Tesla driver/occupant death AND cyclist/pedestrian: {count}")
```

```
Index(['Case #', 'Year', 'Date', 'Country', 'State', 'Description', 'Deaths',
      'Tesla driver', 'Tesla occupant', 'Other vehicle', 'Cyclists/ Peds',
      'TSLA+cycl / peds', 'Model', 'Autopilot claimed',
      'Verified Tesla Autopilot Deaths',
      'Verified Tesla Autopilot Deaths + All Deaths Reported to NHTSA SGO',
      'Deceased 1', 'Deceased 2', 'Deceased 3', 'Deceased 4', 'DayOfWeek'],
      dtype='object')
```

Number of accidents involving Tesla driver/occupant death AND cyclist/pedestrian: 1

- What is the frequency of Tesla colliding with other vehicles ?

```
In [ ]: df['Other vehicle'] = pd.to_numeric(df['Other vehicle'], errors='coerce')
print("Frequency of Tesla colliding with other vehicles:")
count_other_vehicles=df['Other vehicle'].value_counts().sort_index()
print(count_other_vehicles)
```

```
plt.figure(figsize=(8, 5))
count_other_vehicles[count_other_vehicles>=1].plot(kind='bar', color='steelblue')
plt.title('Number of Events by Other Vehicle Collisions')
plt.xlabel('Number of Other Vehicles Involved')
plt.ylabel('Number of Accidents')
plt.xticks(rotation=0)
plt.grid(axis='y')
plt.tight_layout()
plt.show()
```

Frequency of Tesla colliding with other vehicles:

Other vehicle

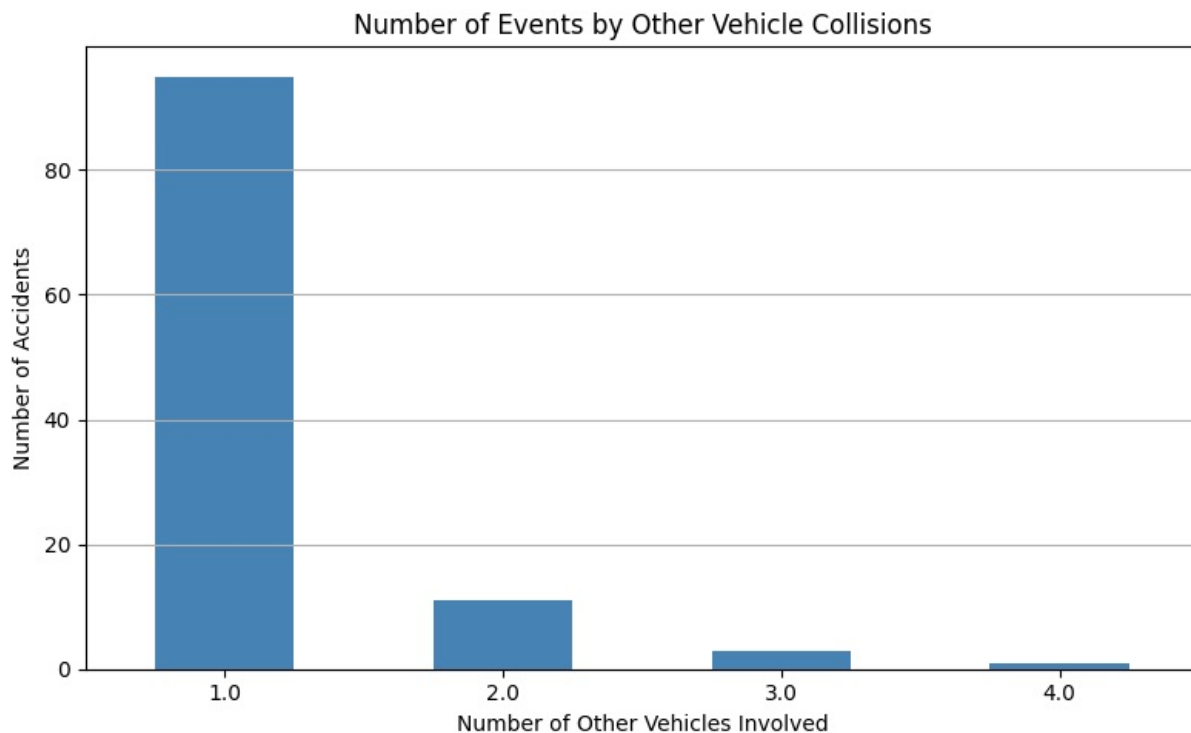
1.0 95

2.0 11

3.0 3

4.0 1

Name: count, dtype: int64



### c. Study the event distribution across models

```
In [ ]: # Standardize formatting
df['Model'] = df['Model'].str.strip().str.title()

# View unique models
print("Unique Tesla Models:")
print(df['Model'].unique())

model_counts = df['Model'].value_counts()

print("Event distribution by Tesla model:")
print(model_counts)

plt.figure(figsize=(8, 5))
model_counts.plot(kind='bar', color='teal')
plt.title('Event Distribution Across Tesla Models')
plt.xlabel('Tesla Model')
plt.ylabel('Number of Events')
plt.xticks(rotation=45)
plt.grid(axis='y')
plt.tight_layout()
plt.show()

model_percentage = model_counts / model_counts.sum() * 100
print("Percentage distribution by model:")
print(model_percentage.round(2))

# Ensure 'Deaths' is numeric
df['Deaths'] = pd.to_numeric(df['Deaths'], errors='coerce')

# Total deaths per model
deaths_per_model = df.groupby('Model')['Deaths'].sum().sort_values(ascending=False)
```

```

print("Total deaths per Tesla model:")
print(deaths_per_model)

deaths_per_model.plot(kind='bar', color='darkred', figsize=(8, 5))
plt.title('Total Deaths by Tesla Model')
plt.xlabel('Tesla Model')
plt.ylabel('Total Deaths')
plt.grid(axis='y')
plt.tight_layout()
plt.show()

df['Autopilot claimed'] = df['Autopilot claimed'].str.strip().str.title()

autopilot_claimed_model = df.groupby(['Model', 'Autopilot claimed']).size().unstack(fill_value=0)

print("Autopilot claimed status by Tesla model:")
print(autopilot_claimed_model)

autopilot_claimed_model.plot(kind='bar', stacked=True, figsize=(9, 6), colormap='Paired')
plt.title('Autopilot Claim Distribution by Tesla Model')
plt.xlabel('Tesla Model')
plt.ylabel('Number of Events')
plt.xticks(rotation=45)
plt.tight_layout()
plt.legend(title='Autopilot Claimed')
plt.grid(axis='y')
plt.show()

df['Cyclists/ Peds'] = pd.to_numeric(df['Cyclists/ Peds'], errors='coerce')
cycl_peds_per_model = df.groupby('Model')['Cyclists/ Peds'].sum().sort_values(ascending=False)

print("Cyclist/Pedestrian hits per Tesla model:")
print(cycl_peds_per_model)
cycl_peds_per_model.plot(kind='bar', color='blue', figsize=(8, 5))
plt.title('Cyclist/Pedestrian hits per Tesla model')
plt.xlabel('Tesla Model')
plt.ylabel('Total Cyclist/Pedestrian hits')
plt.grid(axis='y')
plt.tight_layout()
plt.show()

```

Unique Tesla Models:

[nan 'Y' '1' '2' '3' 'S' 'X']

Event distribution by Tesla model:

Model

S 45

3 39

X 17


Y 10

1 1

2 1

Name: count, dtype: int64

Event Distribution Across Tesla Models



Percentage distribution by model:

Model

S 39.82  
3 34.51  
X 15.04  
Y 8.85  
1 0.88  
2 0.88


Name: count, dtype: float64

Total deaths per Tesla model:

Model

S 53.0  
3 47.0  
X 19.0  
Y 13.0  
2 1.0  
1 1.0

Name: Deaths, dtype: float64

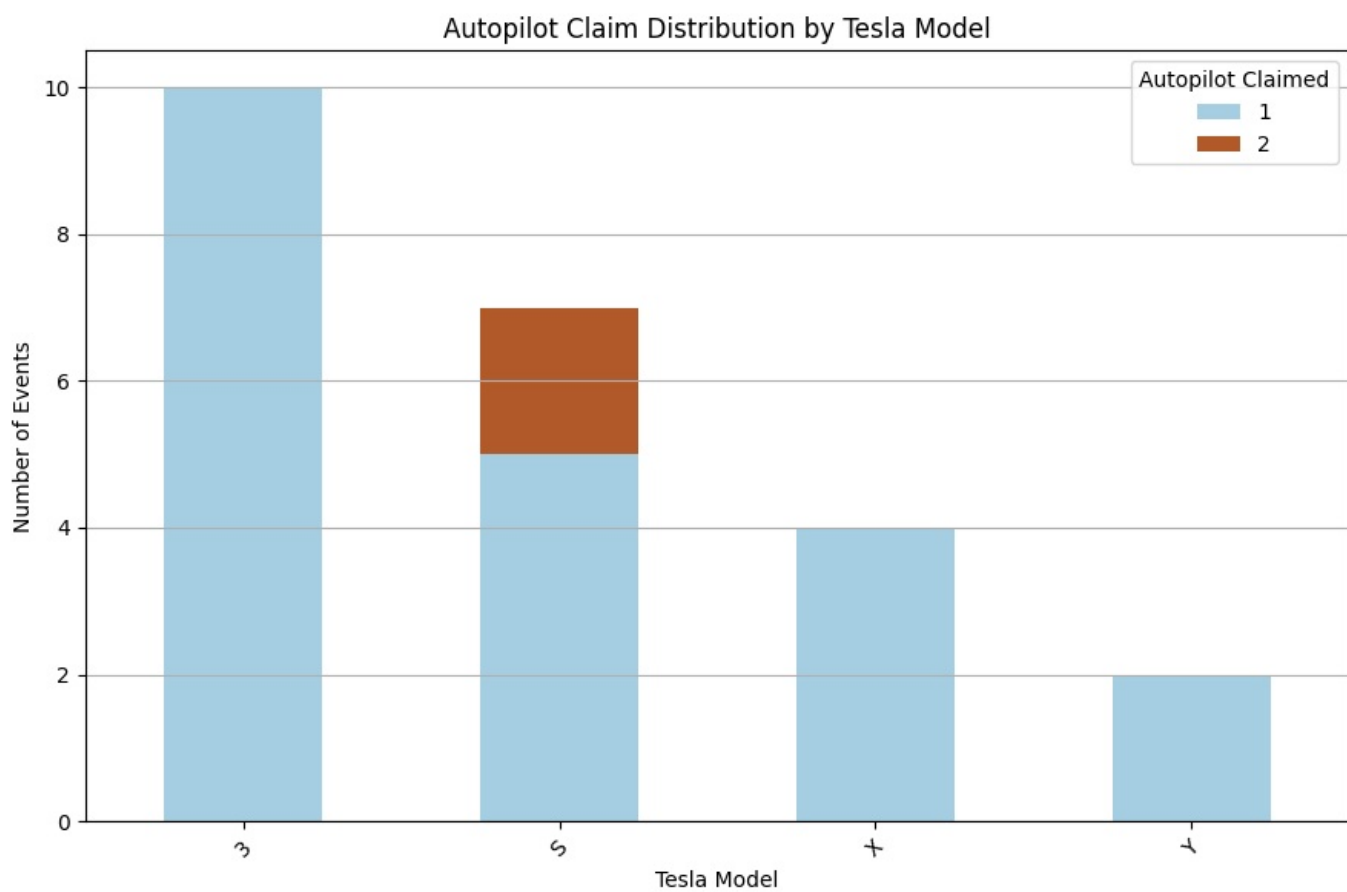


Autopilot claimed status by Tesla model:

Autopilot claimed 1 2

Model

3 10 0  
S 5 2  
X 4 0  
Y 2 0



Cyclist/Pedestrian hits per Tesla model:

Model

3 6.0

S 5.0


X 5.0

Y 2.0

2 0.0

1 0.0

Name: Cyclists/ Peds, dtype: float64



**d. Check the distribution of verified Tesla autopilot deaths**

```
In [ ]: # Ensure it's numeric
df['Verified Tesla Autopilot Deaths'] = pd.to_numeric(df['Verified Tesla Autopilot Deaths'], errors='coerce')
print("Summary of Verified Tesla Autopilot Deaths:")
print(df['Verified Tesla Autopilot Deaths'].describe())
# Count of how many events had 0, 1, 2,... verified deaths
print(df['Verified Tesla Autopilot Deaths'].value_counts().sort_index())

plt.figure(figsize=(7, 5))
df['Verified Tesla Autopilot Deaths'].dropna().astype(int).plot.hist(bins=range(0, int(df['Verified Tesla Autopilot Deaths'].max())+1))
plt.title('Distribution of Verified Tesla Autopilot Deaths')
plt.xlabel('Number of Verified Autopilot Deaths')
plt.ylabel('Number of Events')
plt.grid(axis='y')
plt.tight_layout()
plt.show()

# Assuming 'Year' is available
df['Year'] = pd.to_numeric(df['Year'], errors='coerce')

autopilot_deaths_yearly = df.groupby('Year')['Verified Tesla Autopilot Deaths'].sum()

autopilot_deaths_yearly.plot(kind='line', marker='o', figsize=(8, 5), color='crimson')
plt.title('Yearly Verified Tesla Autopilot Deaths')
plt.xlabel('Year')
plt.ylabel('Number of Verified Deaths')
plt.grid(True)
plt.tight_layout()
plt.show()
```

Summary of Verified Tesla Autopilot Deaths:

count	23.000000
mean	186.652174
std	579.648295
min	1.000000
25%	1.000000
50%	1.000000
75%	9.500000
max	2022.000000

Name: Verified Tesla Autopilot Deaths, dtype: float64

Verified Tesla Autopilot Deaths

1.0	13
2.0	3
3.0	1
16.0	1
19.0	1
75.0	1
118.0	1
2021.0	1
2022.0	1

Name: count, dtype: int64

